# Tech war: Chinese AI models deemed a security risk by new US government report

scmp.com/tech/tech-war/article/3327613/tech-war-chinese-ai-models-deemed-security-risk-new-us-government-report

Vincent Chow October 2, 2025

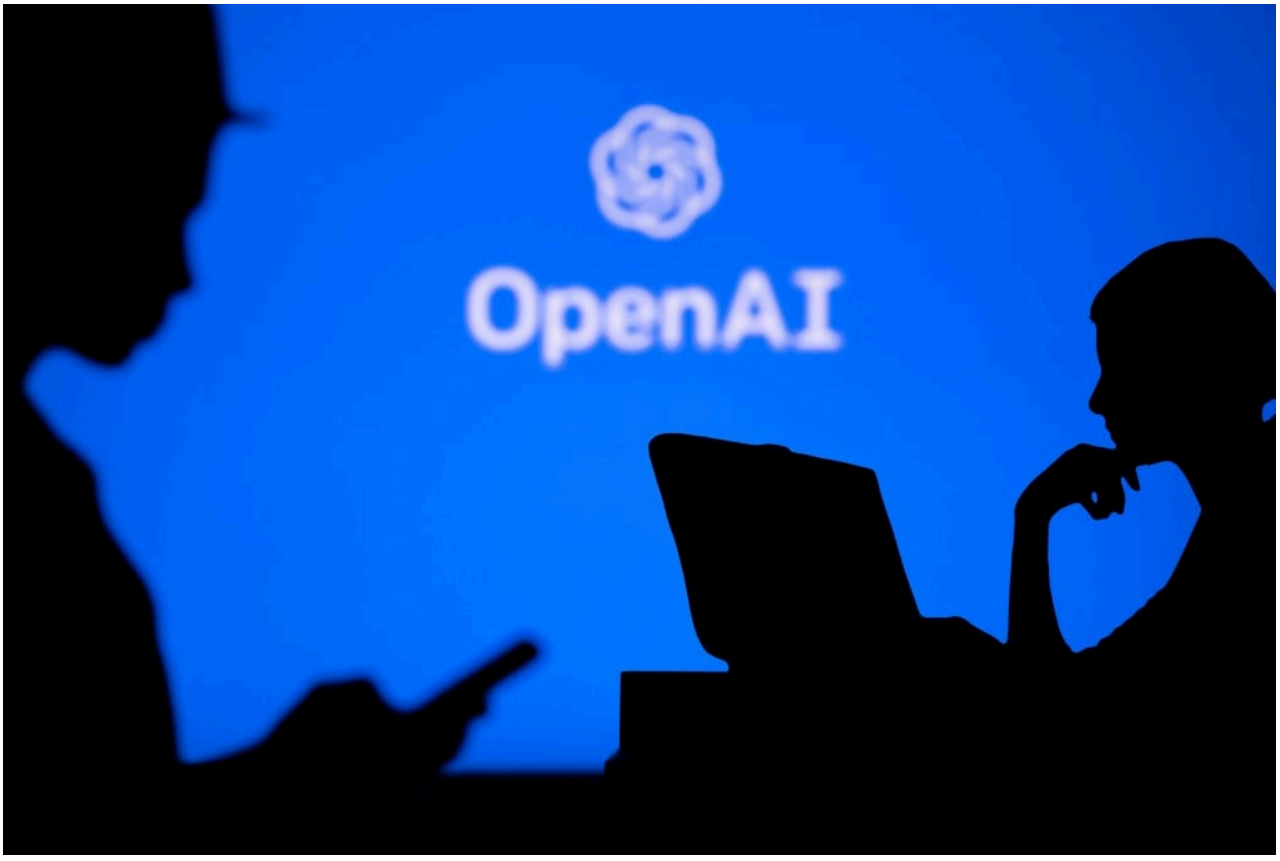


Artificial intelligence
TechTech War

The evaluation marks the first time the US government has comprehensively assessed DeepSeek's capabilities in relation to leading US models

Reading Time:**3 minutes**

Why you can trust SCMP

6

Listen
Vincent Chow
Published: 6:00pm, 2 Oct 2025Updated: 9:45pm, 2 Oct 2025

Chinese models lag behind their American counterparts in performance, cost, security and adoption, despite their growing global popularity, according to a new report from the US government.

Describing Chinese models as “adversary AI”, the report released on Tuesday by the Centre for AI Standards and Innovation (CAISI) at the National Institute of Standards and Technology (NIST) and the Department of Commerce claimed that models such as those from DeepSeek posed risks to AI developers, consumers and US national security because of their security shortcomings and censorship.

The report comes after US President Donald Trump’s AI Action Plan, released in July, called for the evaluation of frontier Chinese models’ capabilities and alignment with state narratives.

DeepSeek, China’s most high-profile AI company, has come under fire in the US where it has been accused of stealing user data and amplifying Chinese state narratives.

According to the report, DeepSeek’s models had lower scores than US models almost across the board on 19 public and internal benchmarks. Photo: Shutterstock Images

The evaluation conducted by CAISI marks the first time the US government has produced a comprehensive assessment of DeepSeek’s capabilities and popularity in relation to leading US models, including OpenAI’s GPT-5, its open-sourced model gpt-oss and Anthropic’s Claude Opus 4.

According to the report, DeepSeek's models had lower scores than US models almost across the board on 19 public and internal benchmarks, while also being more vulnerable to being jailbroken by malicious users intent on carrying out hacking and cybercrime activities.

The report also claimed that Chinese state censorship was "built directly into DeepSeek models", based on a new benchmark CAISI developed in conjunction with the Department of State that tested the models on questions considered politically sensitive to China's ruling Communist Party.

It found that DeepSeek models were more aligned with Chinese state narratives than the US models, with the most aligned model being DeepSeek's R1-0528 at 25.7 per cent when prompted in Chinese.

DeepSeek did not respond to a request for comment.

On the other hand, the Hangzhou-based start-up's "open-weight" models have helped China catch up with the US in the global AI adoption race.

Downloads of DeepSeek models on the developer platform Hugging Face have increased nearly 1,000 per cent since January, while Alibaba Cloud's Qwen family of models increased 135 per cent, the report found.

Alibaba Cloud is the AI and cloud services unit of Alibaba Group Holding, owner of the Post.

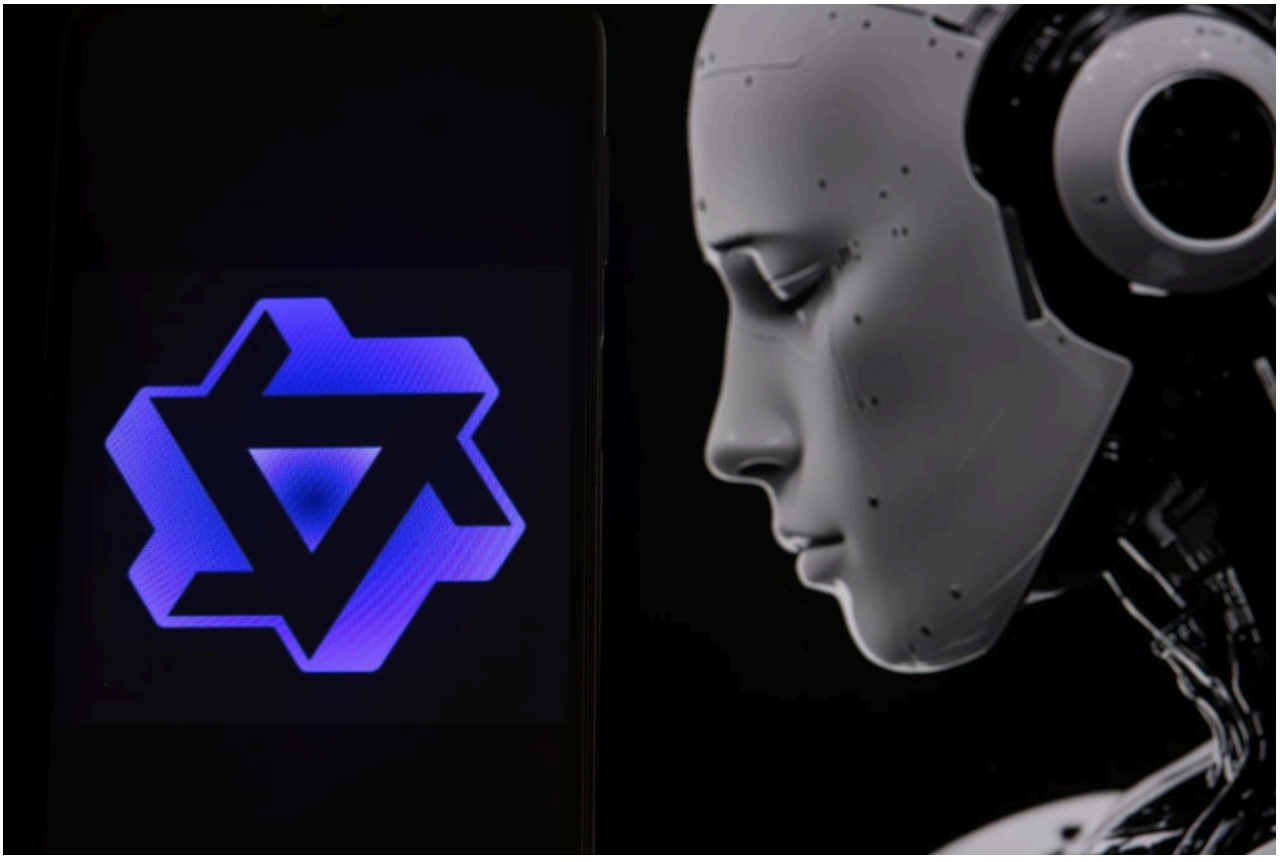

Since January, downloads of Alibaba's Qwen on the developer platform Hugging Face have increased 135 per cent, the report found. Photo: Shutterstock Images

US firms still have the most global downloads across all models on the platform, though Alibaba Cloud is closing in on Meta Platforms, the developer behind the Llama family of models, as the second all-time most popular model maker, only behind OpenAI.

Commonly referred to as open-source, open-weight models are models whose weights – the variables encoding their “intelligence” – are publicly released, allowing developers to download and build on them.

Notably, the number of derivative models built on Qwen and shared on Hugging Face exceeded those from Google, Meta, Microsoft and OpenAI combined.

Another key claim in the report was that OpenAI’s GPT-5-mini cost 35 per cent less on average than DeepSeek’s leading model, V3.1, to perform at a similar level, using their respective application programming interface (API) prices as the basis for comparison.

However, the report did not mention that DeepSeek users have the option of deploying the open-weight models locally, which is not the case for users of proprietary US models who must pay for API access.

DeepSeek has also released newer models in recent weeks, cutting official API prices by over 50 per cent while maintaining similar levels of performance, according to third-party AI benchmarking firm Artificial Analysis.

These updated versions were not included in CAISI’s evaluations, which only looked at DeepSeek’s R1, R1-0528 and V3.1 models.

On social media, US Commerce Secretary Howard Lutnick said his department was helping ensure “continued US leadership in AI” by publishing these findings.

“The report is clear: DeepSeek lags far behind, especially in cyber and software engineering,” Lutnick wrote. “These weaknesses aren’t just technical. They demonstrate why relying on foreign AI is dangerous and shortsighted.”

Launched in 2023 under the Biden administration, the CAISI was previously known as the US AI Safety Institute before having its name changed when Trump came into office with a new focus on promoting US leadership in AI innovation.

The body works directly with US firms such as Anthropic and OpenAI on AI safety and security issues.

Vincent Chow

Vincent Chow is a technology reporter covering AI, with a focus on how society navigates the emergence of increasingly powerful AI systems. He previously covered Chinese society and was awarded a SOPA award for his culture reporting. He is currently supported by the Tarbell Center for AI Journalism.